

Part 1 Data Processing and Engineering

December 20, 2021

1 Predict Future Sales Part 1: Data Processing and Engineering

Motivation

This is the final project for the course “How to win a data science competition”. In this competition you will work with a challenging time-series dataset consisting of daily sales data, kindly provided by one of the largest Russian software firms - 1C Company.

We are asking you to **predict total sales for every product and store in the next month**. By solving this competition you will be able to apply and enhance your data science skills.

You are provided with daily historical sales data. The task is to forecast the total amount of products sold in every shop for the test set. **Note that the list of shops and products slightly changes every month**. Creating a robust model that can handle such situations is part of the challenge.

Source: <https://www.kaggle.com/c/competitive-data-science-predict-future-sales/overview>

File Description - **sales_train.csv** - the training set. Daily historical data from January 2013 to October 2015. - **test.csv** - the test set. You need to forecast the sales for these shops and products for November 2015. - **sample_submission.csv** - a sample submission file in the correct format. - **items.csv** - supplemental information about the items/products. - **item_categories.csv** - supplemental information about the items categories. - **shops.csv**- supplemental information about the shops.

2 Table of contents

- 1. Importing data
- 2. Preprocessing data
 - 2.1 Extracting categorical features from Russian item names
 - 2.2 Creating text features with TF-IDF
 - 2.3 Check duplicates
 - 2.4 Preprocessing for data engineering
 - 2.5 Create lagged features
 - 2.6 Creating time series trend features
 - 2.7 Feature matrix
- 3. Exploratory data analysis
 - 3.1 Target variable
 - 3.2 Multivariate heatmaps
- 4. Advanced Feature Engineering

- 4.1 Mean encoding on categorical features
- 4.2 Matrix factorization of TFIDF processed features

```
[2]: #load packages
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
from math import ceil
import copy

import os
import time
import gc
import pickle
from tqdm import tqdm #progress bar
from itertools import product
import warnings
warnings.filterwarnings("ignore")
from IPython.display import clear_output

from sklearn.feature_extraction import text
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import KFold, StratifiedKFold
from sklearn.metrics import r2_score, mean_squared_error
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA, TruncatedSVD, NMF
from sklearn.linear_model import LinearRegression, Ridge, Lasso, Lars,
↳ ElasticNet
# from sklearn.externals import joblib
import joblib

import lightgbm as lgb
import xgboost as xgb
```

3 1. Importing data

```
[3]: data_folder = "./data/"
sales_train = pd.read_csv(os.path.join(data_folder, "sales_train.csv"))
items = pd.read_csv(os.path.join(data_folder, "items.csv"))
test = pd.read_csv(os.path.join(data_folder, "test.csv")).set_index('ID')
item_categories = pd.read_csv(os.path.join(data_folder, "item_categories.csv"))
shops = pd.read_csv(os.path.join(data_folder, "shops.csv"))
```

```
[11]: sales_train['sales'] = sales_train['item_cnt_day'] * sales_train['item_price']
print('shape: \n', sales_train.shape)
print('-----')
print('dtype: \n', sales_train.dtypes)
print('-----')
print('# of unique values \n', sales_train.nunique())
print('-----')
sales_train.head()
```

shape:

(2935849, 7)

dtype:

date	object
date_block_num	int64
shop_id	int64
item_id	int64
item_price	float64
item_cnt_day	float64
sales	float64

dtype: object

of unique values

date	1034
date_block_num	34
shop_id	60
item_id	21807
item_price	19993
item_cnt_day	198
sales	24771

dtype: int64

```
[11]:
```

	date	date_block_num	shop_id	item_id	item_price	item_cnt_day	\
0	02.01.2013	0	59	22154	999.00	1.0	
1	03.01.2013	0	25	2552	899.00	1.0	
2	05.01.2013	0	25	2552	899.00	-1.0	
3	06.01.2013	0	25	2554	1709.05	1.0	
4	15.01.2013	0	25	2555	1099.00	1.0	

	sales
0	999.00
1	899.00
2	-899.00
3	1709.05
4	1099.00

```
[6]: sales_train.describe()
```

```
[6]:
```

	date_block_num	shop_id	item_id	item_price	item_cnt_day \
count	2.935849e+06	2.935849e+06	2.935849e+06	2.935849e+06	2.935849e+06
mean	1.456991e+01	3.300173e+01	1.019723e+04	8.908532e+02	1.242641e+00
std	9.422988e+00	1.622697e+01	6.324297e+03	1.729800e+03	2.618834e+00
min	0.000000e+00	0.000000e+00	0.000000e+00	-1.000000e+00	-2.200000e+01
25%	7.000000e+00	2.200000e+01	4.476000e+03	2.490000e+02	1.000000e+00
50%	1.400000e+01	3.100000e+01	9.343000e+03	3.990000e+02	1.000000e+00
75%	2.300000e+01	4.700000e+01	1.568400e+04	9.990000e+02	1.000000e+00
max	3.300000e+01	5.900000e+01	2.216900e+04	3.079800e+05	2.169000e+03


```

sales
count 2.935849e+06
mean 1.157732e+03
std 5.683604e+03
min -6.897000e+04
25% 2.490000e+02
50% 4.490000e+02
75% 1.078200e+03
max 1.829990e+06

```

Note The minimum item_price is negative, which is strange.

```
[12]: print('shape: \n', items.shape)
print('-----')
print('# of unique values \n', items.nunique())
print('-----')
items.head()
```

```

shape:
(22170, 3)
-----
# of unique values
item_name      22170
item_id        22170
item_category_id      84
dtype: int64
-----

```

```
[12]:
```

	item_name	item_id \
0	!	0
1	!ABBY FineReader 12 Professional Edition Full...	1
2	*** (UNV)	2
3	*** (Univ)	3
4	*** ()	4


```

item_category_id
0      40
1      76

```

2	40
3	40
4	40

```
[13]: print('shape: \n', test.shape)
print('-----')
print('# of unique values \n', test.nunique())
print('-----')
test.head()
```

```
shape:
(214200, 2)
```

```
-----
# of unique values
shop_id      42
item_id     5100
dtype: int64
-----
```

```
[13]:      shop_id  item_id
ID
0         5      5037
1         5      5320
2         5      5233
3         5      5232
4         5      5268
```

```
[14]: print('shape: \n', item_categories.shape)
print('-----')
print('# of unique values \n', item_categories.nunique())
print('-----')
item_categories.head()
```

```
shape:
(84, 2)
```

```
-----
# of unique values
item_category_name      84
item_category_id       84
dtype: int64
-----
```

```
[14]:      item_category_name  item_category_id
0  PC - /                0
1          - PS2          1
2          - PS3          2
3          - PS4          3
4          - PSP          4
```

```
[15]: print('shape: \n', shops.shape)
print('-----')
print('# of unique values \n', shops.nunique())
print('-----')
shops.head()
```

shape:

(60, 2)

of unique values

shop_name 60

shop_id 60

dtype: int64

```
[15]:      shop_name  shop_id
0      !      , 56      0
1      !      "      "      1
2      "      "      "      2
3      "      -      "      3
4      "      "      "      4
```

4 2. Preprocessing data

In test dataset, each ID represents a unique pair of (shop_id, item_id)

4.1 2.1 Extracting categorical features from Russian item names

```
[16]: shops.head()
```

```
[16]:      shop_name  shop_id
0      !      , 56      0
1      !      "      "      1
2      "      "      "      2
3      "      -      "      3
4      "      "      "      4
```

```
[17]: shops.loc[shops.shop_name == '      "7 "', 'shop_name'] = '      "7 "'
      ↪ #delete space between
shops['city'] = shops['shop_name'].str.split(' ').map(lambda x: x[0])
shops.loc[shops.city == '! ', 'city'] = ' '
shops['city_code'] = LabelEncoder().fit_transform(shops['city'])
shops = shops[['shop_id', 'city_code']]
shops.head()
```

```
[17]:      shop_id  city_code
0          0          29
```

1	1	29
2	2	0
3	3	1
4	4	2

```
[18]: item_categories['split'] = item_categories['item_category_name'].str.split('-')
item_categories['type'] = item_categories['split'].map(lambda x: x[0].strip())
item_categories['type_code'] = LabelEncoder().
    ↪fit_transform(item_categories['type'])
item_categories['subtype'] = item_categories['split'].map(lambda x: x[1].
    ↪strip() if len(x) > 1 else x[0].strip())
item_categories['subtype_code'] = LabelEncoder().
    ↪fit_transform(item_categories['subtype'])
item_categories = item_categories[['item_category_id', 'type_code',
    ↪'subtype_code']]
item_categories.head()
```

```
[18]:   item_category_id  type_code  subtype_code
0              0         0         29
1              1         1          9
2              2         1         10
3              3         1         11
4              4         1         13
```

```
[19]: items.drop('item_name', axis=1, inplace=True)
items.head()
```

```
[19]:   item_id  item_category_id
0         0                40
1         1                76
2         2                40
3         3                40
4         4                40
```

4.2 2.2 Creating text features with TF-IDF

```
[20]: items1= pd.read_csv(os.path.join(data_folder, 'items.csv'))
item_categories1 = pd.read_csv(os.path.join(data_folder, 'item_categories.csv'))
shops1 = pd.read_csv(os.path.join(data_folder, 'shops.csv'))
```

```
[21]: #for items1 names
nb_features = 25
tfidf = text.TfidfVectorizer(max_features=nb_features)
items1['item_name_len'] = items1['item_name'].map(len) #name length
items1['item_name_wc'] = items1['item_name'].map(lambda x: len(str(x).split('
    ↪'))) #word count
txtFeatures = pd.DataFrame(tfidf.fit_transform(items1['item_name']).toarray())
```

```

cols = txtFeatures.columns
for i in range(nb_features):
    items1['item_name_tfidf_'+ str(i)] = txtFeatures[cols[i]]

items1.head()
#txtFeatures.head()

```

```

[21]:

```

	item_name	item_id	\
0	! (.) D	0	
1	!ABBY FineReader 12 Professional Edition Full...	1	
2	*** (UNV) D	2	
3	*** (Univ) D	3	
4	*** () D	4	

	item_category_id	item_name_len	item_name_wc	item_name_tfidf_0	\
0		40	41	14	0.0
1		76	68	9	0.0
2		40	45	26	0.0
3		40	47	26	0.0
4		40	43	25	0.0

	item_name_tfidf_1	item_name_tfidf_2	item_name_tfidf_3	item_name_tfidf_4	\
0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	

	...	item_name_tfidf_15	item_name_tfidf_16	item_name_tfidf_17	\
0	...	0.0	0.000000	0.0	
1	...	0.0	0.403761	0.0	
2	...	0.0	0.000000	0.0	
3	...	0.0	0.000000	0.0	
4	...	0.0	0.000000	0.0	

	item_name_tfidf_18	item_name_tfidf_19	item_name_tfidf_20	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	

	item_name_tfidf_21	item_name_tfidf_22	item_name_tfidf_23	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	

4	0.0	0.0	0.0
---	-----	-----	-----

item_name_tfidf_24	
0	0.000000
1	0.483839
2	0.000000
3	0.000000
4	0.000000

[5 rows x 30 columns]

```
[22]: #for shops1 names
nb_features = 25
tfidf = text.TfidfVectorizer(max_features=nb_features)
shops1['shop_name_len'] = shops1['shop_name'].map(len)
shops1['shop_name_wc'] = shops1['shop_name'].map(lambda x: len(str(x).split(' ')))
txtFeatures = pd.DataFrame(tfidf.fit_transform(shops1['shop_name']).toarray())
cols = txtFeatures.columns
for i in range(nb_features):
    shops1['shop_name_tfidf_' + str(i)] = txtFeatures[cols[i]]
shops1.head()
```

```
[22]:
```

	shop_name	shop_id	shop_name_len	shop_name_wc	\
0	! , 56	0	29	4	
1	! " "	1	29	4	
2	" "	2	16	3	
3	" - "	3	30	3	
4	" "	4	24	4	

	shop_name_tfidf_0	shop_name_tfidf_1	shop_name_tfidf_2	shop_name_tfidf_3	\
0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	
3	0.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	0.0	

	shop_name_tfidf_4	shop_name_tfidf_5	...	shop_name_tfidf_15	\
0	0.0	0.0	...	0.0	
1	0.0	0.0	...	0.0	
2	0.0	0.0	...	0.0	
3	0.0	0.0	...	0.0	
4	0.0	0.0	...	0.0	

	shop_name_tfidf_16	shop_name_tfidf_17	shop_name_tfidf_18	\
0	0.0	0.0	0.0	

1	0.0	0.0	0.0
2	0.0	0.0	0.0
3	0.0	0.0	0.0
4	0.0	0.0	0.0

	shop_name_tfidf_19	shop_name_tfidf_20	shop_name_tfidf_21 \
0	0.0	0.0	0.000000
1	0.0	0.0	0.322815
2	0.0	0.0	0.498580
3	1.0	0.0	0.000000
4	0.0	0.0	0.423972

	shop_name_tfidf_22	shop_name_tfidf_23	shop_name_tfidf_24
0	0.0	0.000000	1.000000
1	0.0	0.689588	0.648274
2	0.0	0.000000	0.000000
3	0.0	0.000000	0.000000
4	0.0	0.000000	0.000000

[5 rows x 29 columns]

4.3 2.3 Check duplicates

```
[24]: #sales_train
cols = ['date', 'date_block_num', 'shop_id', 'item_id', 'item_cnt_day']
print(sales_train.duplicated(subset=cols).value_counts())
sales_train.drop_duplicates(subset=cols, inplace=True)
# sales_train.head()
```

```
False    2935825
True      24
dtype: int64
```

There are 24 duplicated rows in sales_train data.

4.4 2.4 Preprocessing for data engineering

```
[25]: train1 = []
cols = ['date_block_num', 'shop_id', 'item_id']
for block_num in range(34): #range of date_block_num is 0~33,
    ↪max(sales_train['date_block_num'])
    temp = sales_train[sales_train.date_block_num == block_num]
    #iterate all combination of block_num, shop_id, item_id
    train1.append(np.array(list(product([block_num], temp['shop_id'].
    ↪unique(), temp['item_id'].unique()))))

train1 = pd.DataFrame(np.vstack(train1), columns=cols)
train1['date_block_num'] = train1['date_block_num'].astype(np.int8)
```

```

train1['shop_id'] = train1['shop_id'].astype(np.int8)
train1['item_id'] = train1['item_id'].astype(np.int16)
train1.sort_values(cols, inplace=True)
train1.head()

```

```

[25]:      date_block_num  shop_id  item_id
      139255           0         0        19
      141495           0         0        27
      144968           0         0        28
      142661           0         0        29
      138947           0         0        32

```

```

[26]: temp = sales_train.groupby(['date_block_num', 'shop_id', 'item_id']).
      ↪agg({'item_cnt_day': ['sum']})
      temp.columns = ['item_cnt_month']
      temp.reset_index(inplace=True)

      print(temp.shape)
      temp.head()

```

(1609124, 4)

```

[26]:      date_block_num  shop_id  item_id  item_cnt_month
      0           0         0        32             6.0
      1           0         0        33             3.0
      2           0         0        35             1.0
      3           0         0        43             1.0
      4           0         0        51             2.0

```

```

[27]: train1 = pd.merge(train1, temp, on=cols, how='left')
      train1['item_cnt_month'] = (train1['item_cnt_month'].fillna(0).clip(0,20).
      ↪astype(np.float16))

      print(train1.shape)
      train1.head()

```

(10913850, 4)

```

[27]:      date_block_num  shop_id  item_id  item_cnt_month
      0           0         0        19             0.0
      1           0         0        27             0.0
      2           0         0        28             0.0
      3           0         0        29             0.0
      4           0         0        32             6.0

```

We are using sales' results from Jan.2013 to Oct.2015 to predict future sales of Nov.2015, so we add the test set to the training set by adding a 34th date_block_num.

```
[28]: test['date_block_num'] = 34
test['date_block_num'] = test['date_block_num'].astype(np.int8)
test['shop_id'] = test['shop_id'].astype(np.int8)
test['item_id'] = test['item_id'].astype(np.int16)
test.head()

train2 = pd.concat([train1,test],ignore_index=True, sort=False, keys=cols)
train2.fillna(0, inplace=True)
print(train2.shape)

train2.tail()
```

(11128050, 4)

```
[28]:      date_block_num  shop_id  item_id  item_cnt_month
11128045             34       45    18454             0.0
11128046             34       45    16188             0.0
11128047             34       45    15757             0.0
11128048             34       45    19648             0.0
11128049             34       45     969             0.0
```

4.5 2.5 Create lagged features

We want to capture past behavior of several item-shop, item category-shop or item-city pairings (among others):

- *date_avg_item_cnt*
- *date_item_avg_item_cnt*
- *date_shop_avg_item_cnt*
- *date_cat_avg_item_cnt*
- *date_shop_cat_avg_item_cnt*
- *date_shop_subtype_avg_item_cnt*
- *date_city_avg_item_cnt*
- *date_shop_type_avg_item_cnt*
- *date_item_city_avg_item_cnt*
- *date_type_avg_item_cnt*
- *date_subtype_avg_item_cnt*

```
[29]: def lag_feature(df, lags, col):
    tmp = df[['date_block_num', 'shop_id', 'item_id', col]]
    for i in lags:
        shifted = tmp.copy()
        shifted.columns = ['date_block_num', 'shop_id', 'item_id',
        ↪ col+' _lag_'+str(i)]
        shifted['date_block_num'] += i
        df = pd.merge(df, shifted, on=['date_block_num', 'shop_id', 'item_id'],
        ↪ how='left')
    return df
```

```
[30]: shops.head()
```

```
[30]:   shop_id  city_code
0         0         29
1         1         29
2         2          0
3         3          1
4         4          2
```

```
[31]: #merge sales_train, test, shops, items, item_categories
train3 = pd.merge(train2, shops, on=['shop_id'], how='left')
train3 = pd.merge(train3, items, on=['item_id'], how='left')
train3 = pd.merge(train3, item_categories, on=['item_category_id'], how='left')

train3.head()
```

```
[31]:   date_block_num  shop_id  item_id  item_cnt_month  city_code \
0                0         0        19             0.0         29
1                0         0        27             0.0         29
2                0         0        28             0.0         29
3                0         0        29             0.0         29
4                0         0        32             6.0         29

   item_category_id  type_code  subtype_code
0                 40          11             4
1                 19           5            10
2                 30           8            55
3                 23           5            16
4                 40          11             4
```

```
[32]: ##### 1. Create 'date_avg_item_cnt'
temp = train3.groupby(['date_block_num']).agg({'item_cnt_month': ['mean']})
temp.columns = ['date_avg_item_cnt']      ###
temp.reset_index(inplace=True)
train3 = pd.merge(train3, temp, on=['date_block_num'], how='left')
train3 = lag_feature(train3, [1], 'date_avg_item_cnt')
train3.drop(['date_avg_item_cnt'], axis=1, inplace=True)
train3.head()
```

```
[32]:   date_block_num  shop_id  item_id  item_cnt_month  city_code \
0                0         0        19             0.0         29
1                0         0        27             0.0         29
2                0         0        28             0.0         29
3                0         0        29             0.0         29
4                0         0        32             6.0         29

   item_category_id  type_code  subtype_code  date_avg_item_cnt_lag_1
```

0	40	11	4	NaN
1	19	5	10	NaN
2	30	8	55	NaN
3	23	5	16	NaN
4	40	11	4	NaN

```
[33]: # ##### 1. Create 'date_avg_item_cnt'
# temp = train3.groupby(['date_block_num']).agg({'item_cnt_month': ['mean']})
# temp.columns = ['date_avg_item_cnt']          ###
# temp.reset_index(inplace=True)
# train3 = pd.merge(train3, temp, on=['date_block_num'], how='left')
# train3 = lag_feature(train3, [1], 'date_avg_item_cnt')
# train3.drop(['date_avg_item_cnt'], axis=1, inplace=True)

##### 2. Create 'date_item_avg_item_cnt'
temp = train3.groupby(['date_block_num', 'item_id']).agg({'item_cnt_month': ['mean']})
temp.columns = ['date_item_avg_item_cnt']      ###
temp.reset_index(inplace=True)
train3 = pd.merge(train3, temp, on=['date_block_num', 'item_id'], how='left')
train3 = lag_feature(train3, [1, 2, 3, 6, 12], 'date_item_avg_item_cnt')
train3.drop(['date_item_avg_item_cnt'], axis=1, inplace=True)

##### 3. Create 'date_shop_avg_item_cnt'
temp = train3.groupby(['date_block_num', 'shop_id']).agg({'item_cnt_month': ['mean']})
temp.columns = ['date_shop_avg_item_cnt']      ###
temp.reset_index(inplace=True)
train3 = pd.merge(train3, temp, on=['date_block_num', 'shop_id'], how='left')
train3 = lag_feature(train3, [1, 2, 3, 6, 12], 'date_shop_avg_item_cnt')
train3.drop(['date_shop_avg_item_cnt'], axis=1, inplace=True)

##### 4. Create 'date_cat_avg_item_cnt'
temp = train3.groupby(['date_block_num', 'item_category_id']).agg({'item_cnt_month': ['mean']})
temp.columns = ['date_cat_avg_item_cnt']       ###
temp.reset_index(inplace=True)
train3 = pd.merge(train3, temp, on=['date_block_num', 'item_category_id'], how='left')
train3 = lag_feature(train3, [1], 'date_cat_avg_item_cnt')
train3.drop(['date_cat_avg_item_cnt'], axis=1, inplace=True)

##### 5. Create 'date_shop_cat_avg_item_cnt'
temp = train3.groupby(['date_block_num', 'shop_id', 'item_category_id']).agg({'item_cnt_month': ['mean']})
temp.columns = ['date_shop_cat_avg_item_cnt']  ###
temp.reset_index(inplace=True)
```

```

train3 = pd.merge(train3, temp, on=['date_block_num', 'shop_id'],
    ↳ 'item_category_id'], how='left')
train3 = lag_feature(train3, [1], 'date_shop_cat_avg_item_cnt')
train3.drop(['date_shop_cat_avg_item_cnt'], axis=1, inplace=True)

##### 6. Create 'date_shop_type_avg_item_cnt'
temp = train3.groupby(['date_block_num', 'shop_id', 'type_code']).
    ↳agg({'item_cnt_month': ['mean']})
temp.columns = ['date_shop_type_avg_item_cnt']    ###
temp.reset_index(inplace=True)
train3 = pd.merge(train3, temp, on=['date_block_num', 'shop_id', 'type_code'],
    ↳how='left')
train3 = lag_feature(train3, [1], 'date_shop_type_avg_item_cnt')
train3.drop(['date_shop_type_avg_item_cnt'], axis=1, inplace=True)

##### 7. Create 'date_type_avg_item_cnt'
temp = train3.groupby(['date_block_num', 'type_code']).agg({'item_cnt_month':
    ↳ ['mean']})
temp.columns = ['date_type_avg_item_cnt']    ###
temp.reset_index(inplace=True)
train3 = pd.merge(train3, temp, on=['date_block_num', 'type_code'], how='left')
train3 = lag_feature(train3, [1], 'date_type_avg_item_cnt')
train3.drop(['date_type_avg_item_cnt'], axis=1, inplace=True)

##### 8. Create 'date_city_avg_item_cnt'
temp = train3.groupby(['date_block_num', 'city_code']).agg({'item_cnt_month':
    ↳ ['mean']})
temp.columns = ['date_city_avg_item_cnt']    ###
temp.reset_index(inplace=True)
train3 = pd.merge(train3, temp, on=['date_block_num', 'city_code'], how='left')
train3 = lag_feature(train3, [1], 'date_city_avg_item_cnt')
train3.drop(['date_city_avg_item_cnt'], axis=1, inplace=True)

##### 9. Create 'date_shop_subtype_avg_item_cnt'
temp = train3.groupby(['date_block_num', 'shop_id', 'subtype_code']).
    ↳agg({'item_cnt_month': ['mean']})
temp.columns = ['date_shop_subtype_avg_item_cnt']    ###
temp.reset_index(inplace=True)
train3 = pd.merge(train3, temp, on=['date_block_num', 'shop_id'],
    ↳ 'subtype_code'], how='left')
train3 = lag_feature(train3, [1], 'date_shop_subtype_avg_item_cnt')
train3.drop(['date_shop_subtype_avg_item_cnt'], axis=1, inplace=True)

##### 10. Create 'date_item_city_avg_item_cnt'
temp = train3.groupby(['date_block_num', 'item_id', 'city_code']).
    ↳agg({'item_cnt_month': ['mean']})

```

```

temp.columns = ['date_item_city_avg_item_cnt']      ###
temp.reset_index(inplace=True)
train3 = pd.merge(train3, temp, on=['date_block_num', 'item_id', 'city_code'],
    ↪how='left')
train3 = lag_feature(train3, [1], 'date_item_city_avg_item_cnt')
train3.drop(['date_item_city_avg_item_cnt'], axis=1, inplace=True)

##### 11. Create 'date_subtype_avg_item_cnt'
temp = train3.groupby(['date_block_num', 'subtype_code']).agg({'item_cnt_month':
    ↪ ['mean']})
temp.columns = ['date_subtype_avg_item_cnt']      ###
temp.reset_index(inplace=True)
train3 = pd.merge(train3, temp, on=['date_block_num', 'subtype_code'],
    ↪how='left')
train3 = lag_feature(train3, [1], 'date_subtype_avg_item_cnt')
train3.drop(['date_subtype_avg_item_cnt'], axis=1, inplace=True)
train3.head()

```

```

[33]:   date_block_num  shop_id  item_id  item_cnt_month  city_code  \
0              0         0         19              0.0         29
1              0         0         27              0.0         29
2              0         0         28              0.0         29
3              0         0         29              0.0         29
4              0         0         32              6.0         29

   item_category_id  type_code  subtype_code  date_avg_item_cnt_lag_1  \
0              40         11              4              NaN
1              19          5              10              NaN
2              30          8              55              NaN
3              23          5              16              NaN
4              40         11              4              NaN

   date_item_avg_item_cnt_lag_1  ...  date_shop_avg_item_cnt_lag_6  \
0              NaN  ...              NaN
1              NaN  ...              NaN
2              NaN  ...              NaN
3              NaN  ...              NaN
4              NaN  ...              NaN

   date_shop_avg_item_cnt_lag_12  date_cat_avg_item_cnt_lag_1  \
0              NaN              NaN
1              NaN              NaN
2              NaN              NaN
3              NaN              NaN
4              NaN              NaN

   date_shop_cat_avg_item_cnt_lag_1  date_shop_type_avg_item_cnt_lag_1  \

```


0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	NaN
4	NaN	NaN

	date_type_avg_item_cnt_lag_1	date_city_avg_item_cnt_lag_1 \
0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	NaN
4	NaN	NaN

	date_shop_subtype_avg_item_cnt_lag_1	date_item_city_avg_item_cnt_lag_1 \
0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	NaN
4	NaN	NaN

	date_subtype_avg_item_cnt_lag_1
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

[5 rows x 27 columns]

```
[34]: train3.columns
```

```
[34]: Index(['date_block_num', 'shop_id', 'item_id', 'item_cnt_month', 'city_code',
        'item_category_id', 'type_code', 'subtype_code',
        'date_avg_item_cnt_lag_1', 'date_item_avg_item_cnt_lag_1',
        'date_item_avg_item_cnt_lag_2', 'date_item_avg_item_cnt_lag_3',
        'date_item_avg_item_cnt_lag_6', 'date_item_avg_item_cnt_lag_12',
        'date_shop_avg_item_cnt_lag_1', 'date_shop_avg_item_cnt_lag_2',
        'date_shop_avg_item_cnt_lag_3', 'date_shop_avg_item_cnt_lag_6',
        'date_shop_avg_item_cnt_lag_12', 'date_cat_avg_item_cnt_lag_1',
        'date_shop_cat_avg_item_cnt_lag_1', 'date_shop_type_avg_item_cnt_lag_1',
        'date_type_avg_item_cnt_lag_1', 'date_city_avg_item_cnt_lag_1',
        'date_shop_subtype_avg_item_cnt_lag_1',
        'date_item_city_avg_item_cnt_lag_1', 'date_subtype_avg_item_cnt_lag_1'],
        dtype='object')
```

```
[35]: # save data to local disk
output = open('./data_processed/train3.pkl', 'wb')
```

```

pickle.dump(train3, output) #764M

output.close()
# train3.to_csv('./data_processed/train3.csv') #1.3G

```

4.6 2.6 Creating time series trend features

```

[36]: #Trend features for price
temp = sales_train.groupby(['item_id']).agg({'item_price': ['mean']})
temp.columns = ['item_avg_item_price']
temp.reset_index(inplace=True)
train4 = pd.merge(train3, temp, on=['item_id'], how='left')

temp = sales_train.groupby(['date_block_num', 'item_id']).agg({'item_price': ['mean']})
temp.columns = ['date_item_avg_item_price']
temp.reset_index(inplace=True)
train4 = pd.merge(train4, temp, on=['date_block_num', 'item_id'], how='left')

lags=[1,2,3,4,5,6]
train4 = lag_feature(train4, lags, 'date_item_avg_item_price')

train4.head()

```

```

[36]:   date_block_num  shop_id  item_id  item_cnt_month  city_code \
0                0        0       19             0.0         29
1                0        0       27             0.0         29
2                0        0       28             0.0         29
3                0        0       29             0.0         29
4                0        0       32             6.0         29

   item_category_id  type_code  subtype_code  date_avg_item_cnt_lag_1 \
0                40         11             4                 NaN
1                19          5            10                 NaN
2                30          8            55                 NaN
3                23          5            16                 NaN
4                40         11             4                 NaN

   date_item_avg_item_cnt_lag_1  ...  date_item_city_avg_item_cnt_lag_1 \
0                NaN  ...                 NaN
1                NaN  ...                 NaN
2                NaN  ...                 NaN
3                NaN  ...                 NaN
4                NaN  ...                 NaN

   date_subtype_avg_item_cnt_lag_1  item_avg_item_price \
0                NaN                28.000000

```

1	NaN	1461.228571
2	NaN	310.010465
3	NaN	1759.285714
4	NaN	249.629240

	date_item_avg_item_price	date_item_avg_item_price_lag_1 \
0	28.000000	NaN
1	2325.000000	NaN
2	549.000000	NaN
3	2397.500000	NaN
4	338.110349	NaN

	date_item_avg_item_price_lag_2	date_item_avg_item_price_lag_3 \
0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	NaN
4	NaN	NaN

	date_item_avg_item_price_lag_4	date_item_avg_item_price_lag_5 \
0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	NaN
4	NaN	NaN

	date_item_avg_item_price_lag_6
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

[5 rows x 35 columns]

```
[37]: train4.columns
```

```
[37]: Index(['date_block_num', 'shop_id', 'item_id', 'item_cnt_month', 'city_code',
            'item_category_id', 'type_code', 'subtype_code',
            'date_avg_item_cnt_lag_1', 'date_item_avg_item_cnt_lag_1',
            'date_item_avg_item_cnt_lag_2', 'date_item_avg_item_cnt_lag_3',
            'date_item_avg_item_cnt_lag_6', 'date_item_avg_item_cnt_lag_12',
            'date_shop_avg_item_cnt_lag_1', 'date_shop_avg_item_cnt_lag_2',
            'date_shop_avg_item_cnt_lag_3', 'date_shop_avg_item_cnt_lag_6',
            'date_shop_avg_item_cnt_lag_12', 'date_cat_avg_item_cnt_lag_1',
            'date_shop_cat_avg_item_cnt_lag_1', 'date_shop_type_avg_item_cnt_lag_1',
            'date_type_avg_item_cnt_lag_1', 'date_city_avg_item_cnt_lag_1',
```

```

'date_shop_subtype_avg_item_cnt_lag_1',
'date_item_city_avg_item_cnt_lag_1', 'date_subtype_avg_item_cnt_lag_1',
'item_avg_item_price', 'date_item_avg_item_price',
'date_item_avg_item_price_lag_1', 'date_item_avg_item_price_lag_2',
'date_item_avg_item_price_lag_3', 'date_item_avg_item_price_lag_4',
'date_item_avg_item_price_lag_5', 'date_item_avg_item_price_lag_6'],
dtype='object')

```

```

[41]: for i in lags:
        train4['delta_price_lag_'+str(i)] = \
            (train4['date_item_avg_item_price_lag_'+str(i)]-train4['item_avg_item_price'])\
            / train4['item_avg_item_price']

def select_trend(row):
    for i in lags:
        if row['delta_price_lag_'+str(i)]:
            return row['delta_price_lag_'+str(i)]
    return 0

train4['delta_price_lag'] = train4.apply(select_trend, axis=1)
train4['delta_price_lag'].fillna(0, inplace=True)

train4.head()

```

```

[41]:  date_block_num  shop_id  item_id  item_cnt_month  city_code  \
0              0         0         19              0.0         29
1              0         0         27              0.0         29
2              0         0         28              0.0         29
3              0         0         29              0.0         29
4              0         0         32              6.0         29

    item_category_id  type_code  subtype_code  date_avg_item_cnt_lag_1  \
0              40         11         4              NaN
1              19         5         10              NaN
2              30         8         55              NaN
3              23         5         16              NaN
4              40         11         4              NaN

    date_item_avg_item_cnt_lag_1  ...  date_item_avg_item_price_lag_4  \
0              NaN  ...              NaN
1              NaN  ...              NaN
2              NaN  ...              NaN
3              NaN  ...              NaN
4              NaN  ...              NaN

    date_item_avg_item_price_lag_5  date_item_avg_item_price_lag_6  \

```

0		NaN		NaN
1		NaN		NaN
2		NaN		NaN
3		NaN		NaN
4		NaN		NaN

	delta_price_lag_1	delta_price_lag_2	delta_price_lag_3	delta_price_lag_4 \
0	NaN	NaN	NaN	NaN
1	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN

	delta_price_lag_5	delta_price_lag_6	delta_price_lag
0	NaN	NaN	0.0
1	NaN	NaN	0.0
2	NaN	NaN	0.0
3	NaN	NaN	0.0
4	NaN	NaN	0.0

[5 rows x 42 columns]

```
[42]: train4.columns
```

```
[42]: Index(['date_block_num', 'shop_id', 'item_id', 'item_cnt_month', 'city_code',
            'item_category_id', 'type_code', 'subtype_code',
            'date_avg_item_cnt_lag_1', 'date_item_avg_item_cnt_lag_1',
            'date_item_avg_item_cnt_lag_2', 'date_item_avg_item_cnt_lag_3',
            'date_item_avg_item_cnt_lag_6', 'date_item_avg_item_cnt_lag_12',
            'date_shop_avg_item_cnt_lag_1', 'date_shop_avg_item_cnt_lag_2',
            'date_shop_avg_item_cnt_lag_3', 'date_shop_avg_item_cnt_lag_6',
            'date_shop_avg_item_cnt_lag_12', 'date_cat_avg_item_cnt_lag_1',
            'date_shop_cat_avg_item_cnt_lag_1', 'date_shop_type_avg_item_cnt_lag_1',
            'date_type_avg_item_cnt_lag_1', 'date_city_avg_item_cnt_lag_1',
            'date_shop_subtype_avg_item_cnt_lag_1',
            'date_item_city_avg_item_cnt_lag_1', 'date_subtype_avg_item_cnt_lag_1',
            'item_avg_item_price', 'date_item_avg_item_price',
            'date_item_avg_item_price_lag_1', 'date_item_avg_item_price_lag_2',
            'date_item_avg_item_price_lag_3', 'date_item_avg_item_price_lag_4',
            'date_item_avg_item_price_lag_5', 'date_item_avg_item_price_lag_6',
            'delta_price_lag_1', 'delta_price_lag_2', 'delta_price_lag_3',
            'delta_price_lag_4', 'delta_price_lag_5', 'delta_price_lag_6',
            'delta_price_lag'],
            dtype='object')
```

```
[44]: dropped_cols = ['item_avg_item_price', 'date_item_avg_item_price']
      for i in lags:
```

```

dropped_cols += ['date_item_avg_item_price_lag_'+str(i)]
dropped_cols += ['delta_price_lag_'+str(i)]

train4.drop(dropped_cols, axis=1, inplace=True)
train4.head()

```

```

[44]:
  date_block_num  shop_id  item_id  item_cnt_month  city_code  \
0                0        0       19             0.0         29
1                0        0       27             0.0         29
2                0        0       28             0.0         29
3                0        0       29             0.0         29
4                0        0       32             6.0         29

  item_category_id  type_code  subtype_code  date_avg_item_cnt_lag_1  \
0                40         11           4                NaN
1                19          5          10                NaN
2                30          8          55                NaN
3                23          5          16                NaN
4                40         11           4                NaN

  date_item_avg_item_cnt_lag_1  ...  date_shop_avg_item_cnt_lag_12  \
0                NaN  ...                NaN
1                NaN  ...                NaN
2                NaN  ...                NaN
3                NaN  ...                NaN
4                NaN  ...                NaN

  date_cat_avg_item_cnt_lag_1  date_shop_cat_avg_item_cnt_lag_1  \
0                NaN                NaN
1                NaN                NaN
2                NaN                NaN
3                NaN                NaN
4                NaN                NaN

  date_shop_type_avg_item_cnt_lag_1  date_type_avg_item_cnt_lag_1  \
0                NaN                NaN
1                NaN                NaN
2                NaN                NaN
3                NaN                NaN
4                NaN                NaN

  date_city_avg_item_cnt_lag_1  date_shop_subtype_avg_item_cnt_lag_1  \
0                NaN                NaN
1                NaN                NaN
2                NaN                NaN
3                NaN                NaN
4                NaN                NaN

```

	date_item_city_avg_item_cnt_lag_1	date_subtype_avg_item_cnt_lag_1 \
0	NaN	NaN
1	NaN	NaN
2	NaN	NaN
3	NaN	NaN
4	NaN	NaN

	delta_price_lag
0	0.0
1	0.0
2	0.0
3	0.0
4	0.0

[5 rows x 28 columns]

```
[45]: #Trend feature for revenue
temp = sales_train.groupby(['date_block_num', 'shop_id']).agg({'sales':['sum']})
temp.columns = ['date_shop_revenue']
temp.reset_index(inplace=True)
train4 = pd.merge(train4, temp, on=['date_block_num', 'shop_id'], how='left')

temp = train4.groupby(['shop_id']).agg({'date_shop_revenue':['mean']})
temp.columns = ['shop_avg_revenue']
temp.reset_index(inplace=True)
train4 = pd.merge(train4, temp, on=['shop_id'], how='left')

train4['delta_revenue'] = (train4['date_shop_revenue'] -
    ↪ train4['shop_avg_revenue']) / train4['shop_avg_revenue']
train4 = lag_feature(train4, [1], 'delta_revenue')

train4.head()
```

```
[45]:   date_block_num  shop_id  item_id  item_cnt_month  city_code \
0                0        0       19             0.0         29
1                0        0       27             0.0         29
2                0        0       28             0.0         29
3                0        0       29             0.0         29
4                0        0       32             6.0         29

   item_category_id  type_code  subtype_code  date_avg_item_cnt_lag_1 \
0                40         11             4                 NaN
1                19          5            10                 NaN
2                30          8            55                 NaN
3                23          5            16                 NaN
4                40         11             4                 NaN
```

	date_item_avg_item_cnt_lag_1	...	date_type_avg_item_cnt_lag_1	\
0	NaN	...	NaN	
1	NaN	...	NaN	
2	NaN	...	NaN	
3	NaN	...	NaN	
4	NaN	...	NaN	

	date_city_avg_item_cnt_lag_1	date_shop_subtype_avg_item_cnt_lag_1	\
0	NaN	NaN	
1	NaN	NaN	
2	NaN	NaN	
3	NaN	NaN	
4	NaN	NaN	

	date_item_city_avg_item_cnt_lag_1	date_subtype_avg_item_cnt_lag_1	\
0	NaN	NaN	
1	NaN	NaN	
2	NaN	NaN	
3	NaN	NaN	
4	NaN	NaN	

	delta_price_lag	date_shop_revenue	shop_avg_revenue	delta_revenue	\
0	0.0	2966412.0	3.319832e+06	-0.106457	
1	0.0	2966412.0	3.319832e+06	-0.106457	
2	0.0	2966412.0	3.319832e+06	-0.106457	
3	0.0	2966412.0	3.319832e+06	-0.106457	
4	0.0	2966412.0	3.319832e+06	-0.106457	

	delta_revenue_lag_1
0	NaN
1	NaN
2	NaN
3	NaN
4	NaN

[5 rows x 32 columns]

```
[46]: train4.drop(['date_shop_revenue', 'shop_avg_revenue', 'delta_revenue'], axis=1,
    ↪inplace=True)
train4['delta_revenue_lag_1'] = train4['delta_revenue_lag_1'].fillna(0.0)
print(train4.shape)
train4.head()
```

(11128050, 29)


```

[46]:  date_block_num  shop_id  item_id  item_cnt_month  city_code  \
0          0          0        19          0.0          29
1          0          0        27          0.0          29
2          0          0        28          0.0          29
3          0          0        29          0.0          29
4          0          0        32          6.0          29

    item_category_id  type_code  subtype_code  date_avg_item_cnt_lag_1  \
0          40          11          4          NaN
1          19          5          10          NaN
2          30          8          55          NaN
3          23          5          16          NaN
4          40          11          4          NaN

    date_item_avg_item_cnt_lag_1  ...  date_cat_avg_item_cnt_lag_1  \
0          NaN  ...          NaN
1          NaN  ...          NaN
2          NaN  ...          NaN
3          NaN  ...          NaN
4          NaN  ...          NaN

    date_shop_cat_avg_item_cnt_lag_1  date_shop_type_avg_item_cnt_lag_1  \
0          NaN          NaN
1          NaN          NaN
2          NaN          NaN
3          NaN          NaN
4          NaN          NaN

    date_type_avg_item_cnt_lag_1  date_city_avg_item_cnt_lag_1  \
0          NaN          NaN
1          NaN          NaN
2          NaN          NaN
3          NaN          NaN
4          NaN          NaN

    date_shop_subtype_avg_item_cnt_lag_1  date_item_city_avg_item_cnt_lag_1  \
0          NaN          NaN
1          NaN          NaN
2          NaN          NaN
3          NaN          NaN
4          NaN          NaN

    date_subtype_avg_item_cnt_lag_1  delta_price_lag  delta_revenue_lag_1
0          NaN          0.0          0.0
1          NaN          0.0          0.0
2          NaN          0.0          0.0
3          NaN          0.0          0.0

```

4 NaN 0.0 0.0

[5 rows x 29 columns]

```
[48]: train4.columns
```

```
[48]: Index(['date_block_num', 'shop_id', 'item_id', 'item_cnt_month', 'city_code',
        'item_category_id', 'type_code', 'subtype_code',
        'date_avg_item_cnt_lag_1', 'date_item_avg_item_cnt_lag_1',
        'date_item_avg_item_cnt_lag_2', 'date_item_avg_item_cnt_lag_3',
        'date_item_avg_item_cnt_lag_6', 'date_item_avg_item_cnt_lag_12',
        'date_shop_avg_item_cnt_lag_1', 'date_shop_avg_item_cnt_lag_2',
        'date_shop_avg_item_cnt_lag_3', 'date_shop_avg_item_cnt_lag_6',
        'date_shop_avg_item_cnt_lag_12', 'date_cat_avg_item_cnt_lag_1',
        'date_shop_cat_avg_item_cnt_lag_1', 'date_shop_type_avg_item_cnt_lag_1',
        'date_type_avg_item_cnt_lag_1', 'date_city_avg_item_cnt_lag_1',
        'date_shop_subtype_avg_item_cnt_lag_1',
        'date_item_city_avg_item_cnt_lag_1', 'date_subtype_avg_item_cnt_lag_1',
        'delta_price_lag', 'delta_revenue_lag_1'],
        dtype='object')
```

```
[47]: # save train4
output = open('./data_processed/train4.pkl', 'wb')
pickle.dump(train4, output)

output.close()
```

4.7 2.7 Feature matrix

```
[49]: df = train4[train4['date_block_num'] > 11]
print(df.shape)

for col in df.columns:
    if ('_lag_' in col) & (df[col].isnull().any()):
        if ('item_cnt' in col):
            df[col].fillna(0, inplace=True)
```

(6639294, 29)

Adding the processed text feature from section 2.2 to df.

```
[50]: df = pd.merge(df, items1, on='item_id', how='left')
df = pd.merge(df, shops1, on='shop_id', how='left')
df.head()
```

```
[50]:   date_block_num  shop_id  item_id  item_cnt_month  city_code  \
0                12        2        27            0.0          0
1                12        2        30            0.0          0
```

2	12	2	31	0.0	0
3	12	2	32	1.0	0
4	12	2	33	1.0	0

	item_category_id_x	type_code	subtype_code	date_avg_item_cnt_lag_1	\
0	19	5	10	0.411133	
1	40	11	4	0.411133	
2	37	11	1	0.411133	
3	40	11	4	0.411133	
4	37	11	1	0.411133	

	date_item_avg_item_cnt_lag_1	...	shop_name_tfidf_15	shop_name_tfidf_16	\
0	0.086975	...	0.0	0.0	
1	1.021484	...	0.0	0.0	
2	0.543457	...	0.0	0.0	
3	1.934570	...	0.0	0.0	
4	0.913086	...	0.0	0.0	

	shop_name_tfidf_17	shop_name_tfidf_18	shop_name_tfidf_19	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	

	shop_name_tfidf_20	shop_name_tfidf_21	shop_name_tfidf_22	\
0	0.0	0.49858	0.0	
1	0.0	0.49858	0.0	
2	0.0	0.49858	0.0	
3	0.0	0.49858	0.0	
4	0.0	0.49858	0.0	

	shop_name_tfidf_23	shop_name_tfidf_24
0	0.0	0.0
1	0.0	0.0
2	0.0	0.0
3	0.0	0.0
4	0.0	0.0

[5 rows x 86 columns]

```
[51]: df.columns.values
```

```
[51]: array(['date_block_num', 'shop_id', 'item_id', 'item_cnt_month',
        'city_code', 'item_category_id_x', 'type_code', 'subtype_code',
        'date_avg_item_cnt_lag_1', 'date_item_avg_item_cnt_lag_1',
        'date_item_avg_item_cnt_lag_2', 'date_item_avg_item_cnt_lag_3',
```

```

'date_item_avg_item_cnt_lag_6', 'date_item_avg_item_cnt_lag_12',
'date_shop_avg_item_cnt_lag_1', 'date_shop_avg_item_cnt_lag_2',
'date_shop_avg_item_cnt_lag_3', 'date_shop_avg_item_cnt_lag_6',
'date_shop_avg_item_cnt_lag_12', 'date_cat_avg_item_cnt_lag_1',
'date_shop_cat_avg_item_cnt_lag_1',
'date_shop_type_avg_item_cnt_lag_1',
'date_type_avg_item_cnt_lag_1', 'date_city_avg_item_cnt_lag_1',
'date_shop_subtype_avg_item_cnt_lag_1',
'date_item_city_avg_item_cnt_lag_1',
'date_subtype_avg_item_cnt_lag_1', 'delta_price_lag',
'delta_revenue_lag_1', 'item_name', 'item_category_id_y',
'item_name_len', 'item_name_wc', 'item_name_tfidf_0',
'item_name_tfidf_1', 'item_name_tfidf_2', 'item_name_tfidf_3',
'item_name_tfidf_4', 'item_name_tfidf_5', 'item_name_tfidf_6',
'item_name_tfidf_7', 'item_name_tfidf_8', 'item_name_tfidf_9',
'item_name_tfidf_10', 'item_name_tfidf_11', 'item_name_tfidf_12',
'item_name_tfidf_13', 'item_name_tfidf_14', 'item_name_tfidf_15',
'item_name_tfidf_16', 'item_name_tfidf_17', 'item_name_tfidf_18',
'item_name_tfidf_19', 'item_name_tfidf_20', 'item_name_tfidf_21',
'item_name_tfidf_22', 'item_name_tfidf_23', 'item_name_tfidf_24',
'shop_name', 'shop_name_len', 'shop_name_wc', 'shop_name_tfidf_0',
'shop_name_tfidf_1', 'shop_name_tfidf_2', 'shop_name_tfidf_3',
'shop_name_tfidf_4', 'shop_name_tfidf_5', 'shop_name_tfidf_6',
'shop_name_tfidf_7', 'shop_name_tfidf_8', 'shop_name_tfidf_9',
'shop_name_tfidf_10', 'shop_name_tfidf_11', 'shop_name_tfidf_12',
'shop_name_tfidf_13', 'shop_name_tfidf_14', 'shop_name_tfidf_15',
'shop_name_tfidf_16', 'shop_name_tfidf_17', 'shop_name_tfidf_18',
'shop_name_tfidf_19', 'shop_name_tfidf_20', 'shop_name_tfidf_21',
'shop_name_tfidf_22', 'shop_name_tfidf_23', 'shop_name_tfidf_24'],
dtype=object)

```

```

[52]: df.drop(['item_name', 'item_name_len', 'item_name_wc', 'item_category_id_y',
              'shop_name', 'shop_name_len', 'shop_name_wc'], axis=1, inplace=True)
print(df.shape)

```

```

(6639294, 79)

```

```

[53]: cat_features = ['shop_id', 'item_id', 'city_code', 'item_category_id_x',
                    ↪ 'type_code', 'subtype_code']

```

```

[54]: del train1, train2, train3, train4, temp
gc.collect()

```

```

[54]: 6753

```

```

[56]: df['item_cnt_month'].unique()

```

```
[56]: array([ 0.,  1.,  2.,  7.,  9.,  3.,  4., 10.,  6.,  5., 20., 13., 11.,
        12.,  8., 16., 18., 19., 14., 15., 17.], dtype=float16)
```

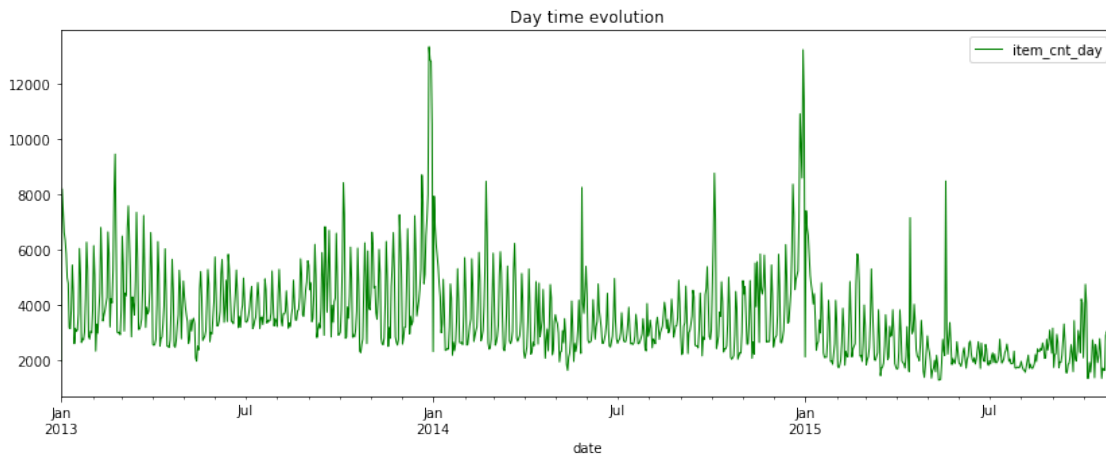
5 3. Exploratory data analysis

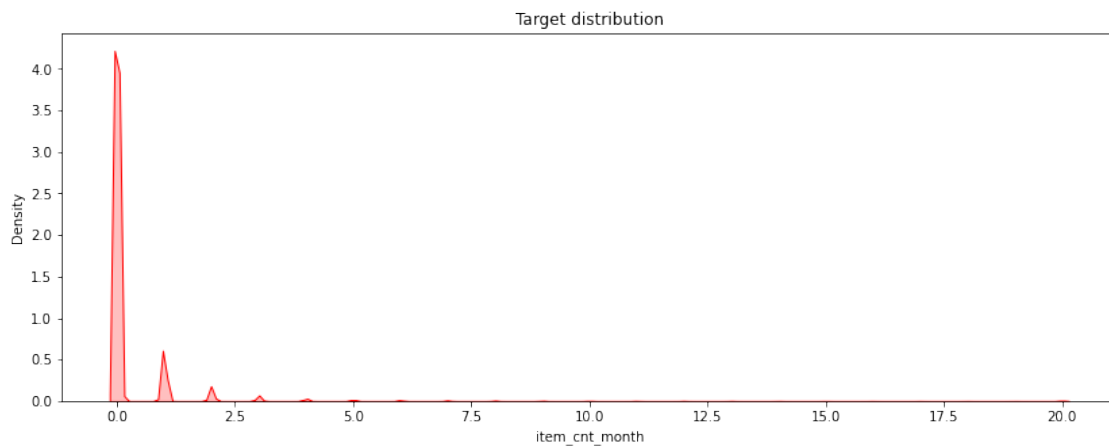
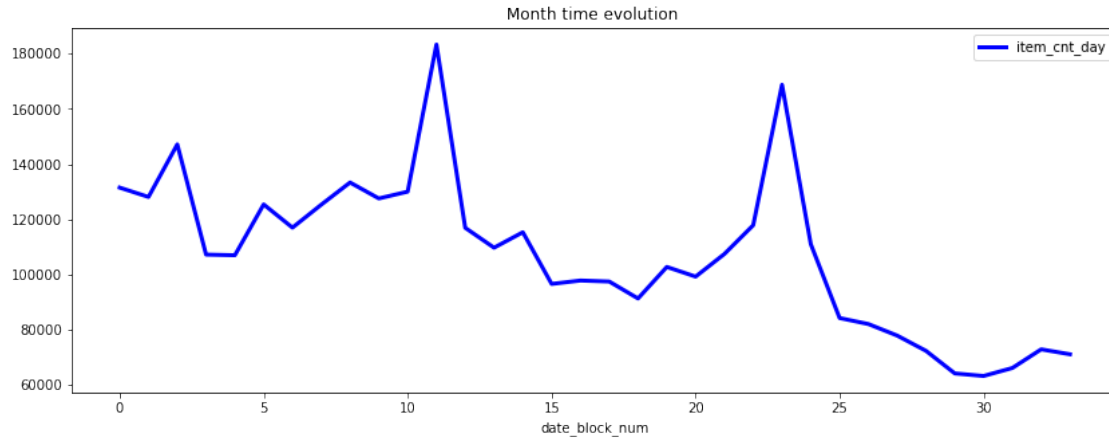
5.1 3.1 Target variable

The **target variable** to be predicted is `item_cnt_month`. Let's check its time series evolution and KDE distribution.

```
[57]: sales_train['date'] = pd.to_datetime(sales_train['date'], format='%d.%m.%Y')
sales_train.groupby('date').agg({'item_cnt_day': 'sum'}).plot(figsize=(14,5),
    ↳title='Day time evolution', lw=1.0, color='green')
sales_train.groupby('date_block_num').agg({'item_cnt_day': 'sum'}).
    ↳plot(figsize=(14,5), title='Month time evolution ', lw=3.0, color='blue')
plt.show()

fig, ax = plt.subplots(figsize=(14,5))
ax = sns.kdeplot(df['item_cnt_month'].values, shade=True, color="red", ax=ax)
ax.set_xlabel('item_cnt_month')
ax.set_ylabel('Density')
ax.set_title('Target distribution')
plt.show()
```





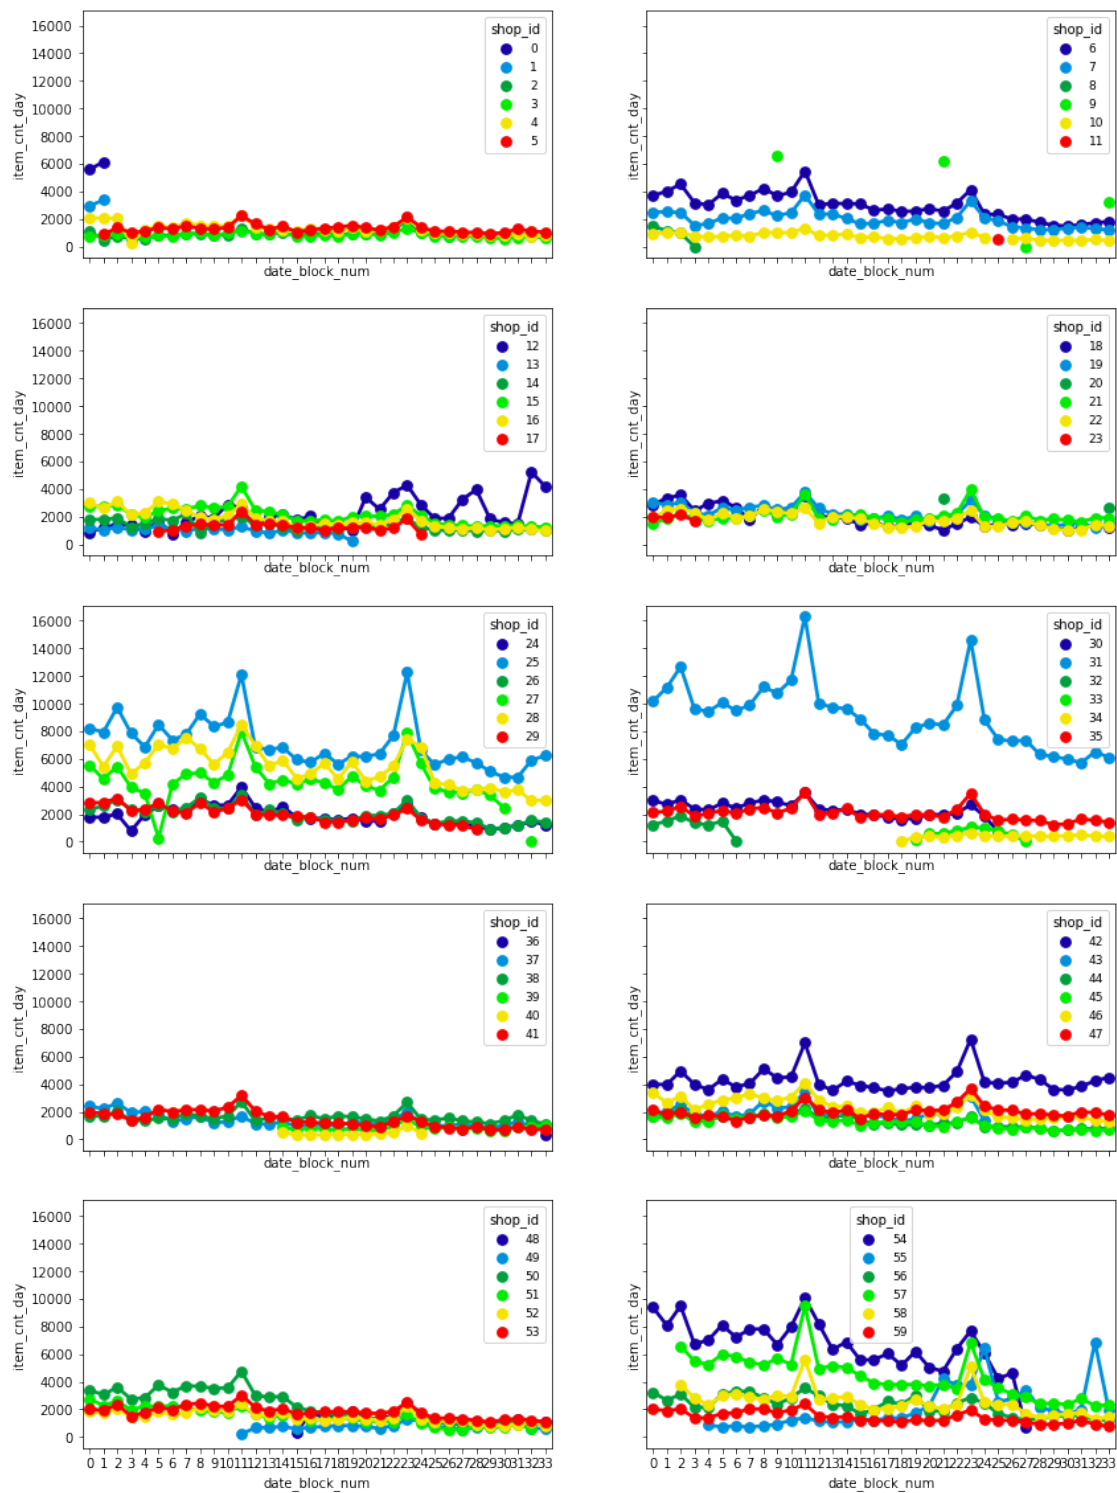
Now check the monthly performance of each shop:

```
[58]: temp = pd.DataFrame(sales_train.groupby(['shop_id',
    ↳ 'date_block_num'])['item_cnt_day'].sum().reset_index())
fig, axes = plt.subplots(nrows=5, ncols=2, sharex=True, sharey=True,
    ↳ figsize=(14,20))
num_graph = 10
id_per_graph = ceil(temp.shop_id.max() / num_graph) #temp.shop_id.max() is 59
x = 0
for i in range(5):
    for j in range(2):
        sns.pointplot(x='date_block_num', y='item_cnt_day', hue='shop_id',
    ↳ palette='nipy_spectral',
                        data = temp[np.logical_and(x * id_per_graph <=
    ↳ temp['shop_id'],
```

```

temp['shop_id'] < (x+1) *
→id_per_graph)],
    ax=axes[i][j], grid=True, lw=1.0)
plt.setp(axes[i][j].get_legend().get_texts(), fontsize='9')
plt.setp(axes[i][j].get_legend().get_title(), fontsize='10')
    x += 1
# print(temp.shop_id.max())
# temp.head(10)

```



5.2 3.2 Multivariate heatmaps

We can pair different categorical variables (ex. `item_category_id`, `shop_id`, `city_code`) together and check the performance of different pairs.


```
[59]: #pair (item_category_id, shop_id)
stores_hm = df.pivot_table(index='shop_id', columns='item_category_id_x',
    ↪values='item_cnt_month', aggfunc='count', fill_value=0)
stores_hm.head()
```

```
[59]: item_category_id_x  0    2    3    4    5    6    7    8    9   11  ...  74  \
shop_id
2                1  513  467  58  265  414  276  49  153  198  ...  22
3                1  513  467  58  265  414  276  49  153  198  ...  22
4                1  513  467  58  265  414  276  49  153  198  ...  22
5                1  513  467  58  265  414  276  49  153  198  ...  22
6                1  513  467  58  265  414  276  49  153  198  ...  22


item_category_id_x   75    76    77    78   79   80   81   82   83
shop_id
2                1368  1290  315  1236   23   36   12   22   92
3                1368  1290  315  1236   23   36   12   22   92
4                1368  1290  315  1236   23   36   12   22   92
5                1368  1290  315  1236   23   36   12   22   92
6                1368  1290  315  1236   23   36   12   22   92
```

[5 rows x 80 columns]

```
[60]: fig, ax = plt.subplots(figsize=(16, 8))
sns.heatmap(stores_hm, cmap='nipy_spectral', ax=ax, cbar=True,
    ↪cbar_kws={'label': 'item_cnt_month'});
```

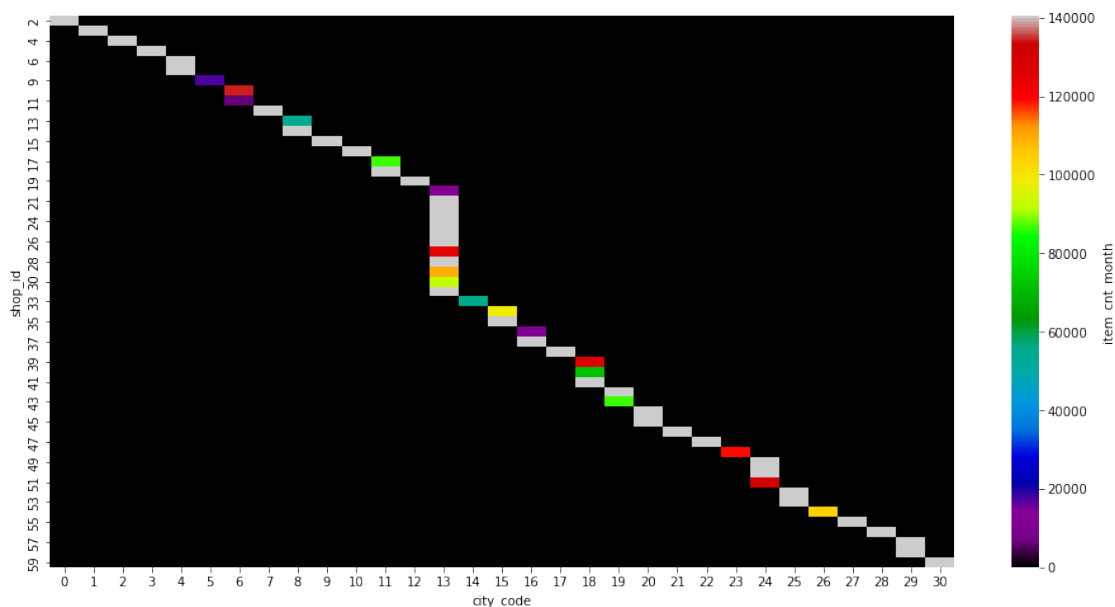


```
[61]: #pair (item_category_id, city_code)
stores_hm = df.pivot_table(index='city_code', columns='item_category_id_x',
    ↪values='item_cnt_month', aggfunc='count', fill_value=0)
# stores_hm.head()
fig, ax = plt.subplots(figsize=(16, 8))
sns.heatmap(stores_hm, cmap='nipy_spectral', ax=ax, cbar=True,
    ↪cbar_kws={'label': 'item_cnt_month'});
```



The next figure reveals the number of shops per cities and their sales performance.

```
[62]: #pair (shop_id, city_code)
stores_hm = df.pivot_table(index='shop_id', columns='city_code',
    ↪ values='item_cnt_month', aggfunc='count', fill_value=0)
# stores_hm.head()
fig, ax = plt.subplots(figsize=(16, 8))
sns.heatmap(stores_hm, cmap='nipy_spectral', ax=ax, cbar=True,
    ↪ cbar_kws={'label': 'item_cnt_month'});
```



Several pairs of shop-item perform better than others, let's exploit this by using target/mean encodings.

6 4. Advancedd Feature Engineering

6.1 4.1 Mean encoding on categorical features

Mean encodings look to map a high cardinality categorical feature into a 1D array (instead of high numbers of them had we used one-hot encoding) based on how often the target variable appears on average in the categorical feature. It is a convenient approach since it also has a clever way of imputing missing values among categorical variables.

```
[45]: import copy
```

```
[63]: df1 = copy.copy(df)
df1.reset_index(inplace=True)

cv = KFold(n_splits=5, shuffle=False)
```

```

new_features = []

check = False
for train_idx, valid_idx in cv.split(df1):
    # Train/validation split
    X_train, X_valid = df1.iloc[train_idx:], df1.iloc[valid_idx:]

    #Mean encoding
    #cat_features = ['shop_id', 'item_id', 'city_code',
    #'item_category_id_x', 'type_code', 'subtype_code']
    for col in cat_features:
        means = X_valid[col].map(X_train.groupby(col).item_cnt_month.mean()) #
        col_new = col + '_target_enc'
        X_valid[col_new] = means

    #results
    df1.loc[valid_idx, col_new] = X_valid

    #store new columns
    if check is False:
        new_features.append(col_new)
    check = True

print(new_features)
# df1.head()

```

```

['shop_id_target_enc', 'item_id_target_enc', 'city_code_target_enc',
'item_category_id_x_target_enc', 'type_code_target_enc',
'subtype_code_target_enc']

```

```
[64]: X_train.head()
```

```

[64]:   index  date_block_num  shop_id  item_id  item_cnt_month  city_code  \
0      0              12        2      27              0.0          0
1      1              12        2      30              0.0          0
2      2              12        2      31              0.0          0
3      3              12        2      32              1.0          0
4      4              12        2      33              1.0          0

      item_category_id_x  type_code  subtype_code  date_avg_item_cnt_lag_1  ...  \
0                    19          5           10          0.411133  ...
1                    40         11            4          0.411133  ...
2                    37         11            1          0.411133  ...
3                    40         11            4          0.411133  ...
4                    37         11            1          0.411133  ...

      shop_name_tfidf_21  shop_name_tfidf_22  shop_name_tfidf_23  \
0          0.49858          0.0          0.0

```

1	0.49858	0.0	0.0
2	0.49858	0.0	0.0
3	0.49858	0.0	0.0
4	0.49858	0.0	0.0

	shop_name_tfidf_24	shop_id_target_enc	item_id_target_enc \
0	0.0	0.141602	0.040802
1	0.0	0.141602	0.186646
2	0.0	0.141602	0.315186
3	0.0	0.141602	0.575195
4	0.0	0.141602	0.336426

	city_code_target_enc	item_category_id_x_target_enc	type_code_target_enc \
0	0.141602	0.550781	0.570312
1	0.141602	0.230835	0.197021
2	0.141602	0.157837	0.197021
3	0.141602	0.230835	0.197021
4	0.141602	0.157837	0.197021

	subtype_code_target_enc
0	0.511230
1	0.230835
2	0.158691
3	0.230835
4	0.158691

[5 rows x 86 columns]

```
[65]: X_valid.head()
```

```
[65]:
```

	index	date_block_num	shop_id	item_id	item_cnt_month	city_code \
5311436	5311436	29	4	2270	0.0	2
5311437	5311437	29	4	2271	0.0	2
5311438	5311438	29	4	2272	0.0	2
5311439	5311439	29	4	2275	0.0	2
5311440	5311440	29	4	2277	0.0	2

	item_category_id_x	type_code	subtype_code	date_avg_item_cnt_lag_1 \
5311436	23	5	16	0.259521
5311437	24	5	17	0.259521
5311438	31	8	61	0.259521
5311439	31	8	61	0.000000
5311440	31	8	61	0.000000

	...	shop_name_tfidf_21	shop_name_tfidf_22	shop_name_tfidf_23 \
5311436	...	0.423972	0.0	0.0
5311437	...	0.423972	0.0	0.0

5311438	...	0.423972	0.0	0.0
5311439	...	0.423972	0.0	0.0
5311440	...	0.423972	0.0	0.0

	shop_name_tfidf_24	shop_id_target_enc	item_id_target_enc	\
5311436	0.0	0.173706	0.752930	
5311437	0.0	0.173706	0.167480	
5311438	0.0	0.173706	0.037903	
5311439	0.0	0.173706	0.032013	
5311440	0.0	0.173706	0.030106	

	city_code_target_enc	item_category_id_x_target_enc	\
5311436	0.173706	0.634766	
5311437	0.173706	0.480469	
5311438	0.173706	0.062103	
5311439	0.173706	0.062103	
5311440	0.173706	0.062103	

	type_code_target_enc	subtype_code_target_enc
5311436	0.633789	0.615723
5311437	0.633789	0.515625
5311438	0.543457	0.070801
5311439	0.543457	0.070801
5311440	0.543457	0.070801

[5 rows x 86 columns]

```
[66]: #fill missing value with mean of item_cnt_month
new_features = ['shop_id_target_enc', 'item_id_target_enc',
               ↪ 'city_code_target_enc', 'item_category_id_x_target_enc',
               ↪ 'type_code_target_enc', 'subtype_code_target_enc']
prior = np.mean(df1['item_cnt_month'].values)
df1[new_features] = df1[new_features].fillna(prior)
df1.head()
```

```
[66]:   index  date_block_num  shop_id  item_id  item_cnt_month  city_code  \
0      0              12        2      27              0.0          0
1      1              12        2      30              0.0          0
2      2              12        2      31              0.0          0
3      3              12        2      32              1.0          0
4      4              12        2      33              1.0          0

   item_category_id_x  type_code  subtype_code  date_avg_item_cnt_lag_1  ...  \
0              19          5          10          0.411133  ...
1              40         11           4          0.411133  ...
2              37         11           1          0.411133  ...
3              40         11           4          0.411133  ...
```

4	37	11	1	0.411133 ...
---	----	----	---	--------------

	shop_name_tfidf_21	shop_name_tfidf_22	shop_name_tfidf_23	\
0	0.49858	0.0	0.0	
1	0.49858	0.0	0.0	
2	0.49858	0.0	0.0	
3	0.49858	0.0	0.0	
4	0.49858	0.0	0.0	

	shop_name_tfidf_24	shop_id_target_enc	item_id_target_enc	\
0	0.0	0.141602	0.040802	
1	0.0	0.141602	0.186646	
2	0.0	0.141602	0.315186	
3	0.0	0.141602	0.575195	
4	0.0	0.141602	0.336426	

	city_code_target_enc	item_category_id_x_target_enc	type_code_target_enc	\
0	0.141602	0.550781	0.570312	
1	0.141602	0.230835	0.197021	
2	0.141602	0.157837	0.197021	
3	0.141602	0.230835	0.197021	
4	0.141602	0.157837	0.197021	

	subtype_code_target_enc
0	0.511230
1	0.230835
2	0.158691
3	0.230835
4	0.158691

[5 rows x 86 columns]

```
[68]: df1.drop(cat_features, axis=1, inplace=True)
df1.drop('index', axis=1, inplace=True)
cat_features = new_features
df1.columns.values
```

```
[68]: array(['date_block_num', 'item_cnt_month', 'date_avg_item_cnt_lag_1',
'date_item_avg_item_cnt_lag_1', 'date_item_avg_item_cnt_lag_2',
'date_item_avg_item_cnt_lag_3', 'date_item_avg_item_cnt_lag_6',
'date_item_avg_item_cnt_lag_12', 'date_shop_avg_item_cnt_lag_1',
'date_shop_avg_item_cnt_lag_2', 'date_shop_avg_item_cnt_lag_3',
'date_shop_avg_item_cnt_lag_6', 'date_shop_avg_item_cnt_lag_12',
'date_cat_avg_item_cnt_lag_1', 'date_shop_cat_avg_item_cnt_lag_1',
'date_shop_type_avg_item_cnt_lag_1',
'date_type_avg_item_cnt_lag_1', 'date_city_avg_item_cnt_lag_1',
'date_shop_subtype_avg_item_cnt_lag_1',
```

```

'date_item_city_avg_item_cnt_lag_1',
'date_subtype_avg_item_cnt_lag_1', 'delta_price_lag',
'delta_revenue_lag_1', 'item_name_tfidf_0', 'item_name_tfidf_1',
'item_name_tfidf_2', 'item_name_tfidf_3', 'item_name_tfidf_4',
'item_name_tfidf_5', 'item_name_tfidf_6', 'item_name_tfidf_7',
'item_name_tfidf_8', 'item_name_tfidf_9', 'item_name_tfidf_10',
'item_name_tfidf_11', 'item_name_tfidf_12', 'item_name_tfidf_13',
'item_name_tfidf_14', 'item_name_tfidf_15', 'item_name_tfidf_16',
'item_name_tfidf_17', 'item_name_tfidf_18', 'item_name_tfidf_19',
'item_name_tfidf_20', 'item_name_tfidf_21', 'item_name_tfidf_22',
'item_name_tfidf_23', 'item_name_tfidf_24', 'shop_name_tfidf_0',
'shop_name_tfidf_1', 'shop_name_tfidf_2', 'shop_name_tfidf_3',
'shop_name_tfidf_4', 'shop_name_tfidf_5', 'shop_name_tfidf_6',
'shop_name_tfidf_7', 'shop_name_tfidf_8', 'shop_name_tfidf_9',
'shop_name_tfidf_10', 'shop_name_tfidf_11', 'shop_name_tfidf_12',
'shop_name_tfidf_13', 'shop_name_tfidf_14', 'shop_name_tfidf_15',
'shop_name_tfidf_16', 'shop_name_tfidf_17', 'shop_name_tfidf_18',
'shop_name_tfidf_19', 'shop_name_tfidf_20', 'shop_name_tfidf_21',
'shop_name_tfidf_22', 'shop_name_tfidf_23', 'shop_name_tfidf_24',
'shop_id_target_enc', 'item_id_target_enc', 'city_code_target_enc',
'item_category_id_x_target_enc', 'type_code_target_enc',
'subtype_code_target_enc'], dtype=object)

```

6.2 4.2 Matrix factorization of TFIDF processed features

We look to reduce the *dimensionality* of the TFIDF-processed features (from 50 dimensions to 10) and extract nonlinear relationships between the text features with **Non-Negative Matrix Factorization (NMF)**:

```

[69]: tfidf_features = ['item_name_tfidf_0', 'item_name_tfidf_1', 'item_name_tfidf_2',
                        'item_name_tfidf_3', 'item_name_tfidf_4', 'item_name_tfidf_5',
                        'item_name_tfidf_6', 'item_name_tfidf_7', 'item_name_tfidf_8',
                        'item_name_tfidf_9', 'item_name_tfidf_10',
                        ↪ 'item_name_tfidf_11',
                        'item_name_tfidf_12', 'item_name_tfidf_13',
                        ↪ 'item_name_tfidf_14',
                        'item_name_tfidf_15', 'item_name_tfidf_16',
                        ↪ 'item_name_tfidf_17',
                        'item_name_tfidf_18', 'item_name_tfidf_19',
                        ↪ 'item_name_tfidf_20',
                        'item_name_tfidf_21', 'item_name_tfidf_22',
                        ↪ 'item_name_tfidf_23',
                        'item_name_tfidf_24', 'shop_name_tfidf_0',
                        ↪ 'shop_name_tfidf_1',
                        'shop_name_tfidf_2', 'shop_name_tfidf_3', 'shop_name_tfidf_4',
                        'shop_name_tfidf_5', 'shop_name_tfidf_6', 'shop_name_tfidf_7',

```



```

        'shop_name_tfidf_8', 'shop_name_tfidf_9',
        ↪ 'shop_name_tfidf_10',
        'shop_name_tfidf_11', 'shop_name_tfidf_12',
        ↪ 'shop_name_tfidf_13',
        'shop_name_tfidf_14', 'shop_name_tfidf_15',
        ↪ 'shop_name_tfidf_16',
        'shop_name_tfidf_17', 'shop_name_tfidf_18',
        ↪ 'shop_name_tfidf_19',
        'shop_name_tfidf_20', 'shop_name_tfidf_21',
        ↪ 'shop_name_tfidf_22',
        'shop_name_tfidf_23', 'shop_name_tfidf_24']

```

```

X_train = df1[df1['date_block_num'] < 34][tfidf_features]
X_test = df1[df1['date_block_num'] == 34][tfidf_features]
print(X_train.shape, X_test.shape)

```

(6425094, 50) (214200, 50)

```

[70]: nmf = NMF(n_components=10, init=None, solver='cd', beta_loss='frobenius', tol=0.
        ↪ 0001, max_iter=200)
nmf.fit(df1[tfidf_features])

```

[70]: NMF(init=None, n_components=10)

```

[71]: X_train = nmf.transform(X_train)
X_test = nmf.transform(X_test)
print(X_train.shape, X_test.shape)

```

(6425094, 10) (214200, 10)

```

[72]: df1.drop(tfidf_features, axis=1, inplace=True)

```

```

[73]: df1.head()

```

```

[73]:   date_block_num  item_cnt_month  date_avg_item_cnt_lag_1 \
0                12              0.0              0.411133
1                12              0.0              0.411133
2                12              0.0              0.411133
3                12              1.0              0.411133
4                12              1.0              0.411133

      date_item_avg_item_cnt_lag_1  date_item_avg_item_cnt_lag_2 \
0                0.086975                0.044434
1                1.021484                1.022461
2                0.543457                0.600098
3                1.934570                1.799805
4                0.913086                0.333252

```

	date_item_avg_item_cnt_lag_3	date_item_avg_item_cnt_lag_6 \
0	0.130493	0.065247
1	0.521973	0.891113
2	0.543457	0.304443
3	1.260742	1.891602
4	0.717285	1.000000

	date_item_avg_item_cnt_lag_12	date_shop_avg_item_cnt_lag_1 \
0	0.155518	0.148071
1	0.000000	0.148071
2	0.000000	0.148071
3	5.378906	0.148071
4	1.355469	0.148071

	date_shop_avg_item_cnt_lag_2 ...	date_item_city_avg_item_cnt_lag_1 \
0	0.100647 ...	0.0
1	0.100647 ...	0.0
2	0.100647 ...	0.0
3	0.100647 ...	0.0
4	0.100647 ...	1.0

	date_subtype_avg_item_cnt_lag_1	delta_price_lag	delta_revenue_lag_1 \
0	1.075195	-0.282795	1.230799
1	0.291504	-0.483264	1.230799
2	0.234009	-0.137618	1.230799
3	0.291504	-0.407143	1.230799
4	0.234009	-0.225177	1.230799

	shop_id_target_enc	item_id_target_enc	city_code_target_enc \
0	0.141602	0.040802	0.141602
1	0.141602	0.186646	0.141602
2	0.141602	0.315186	0.141602
3	0.141602	0.575195	0.141602
4	0.141602	0.336426	0.141602

	item_category_id_x_target_enc	type_code_target_enc \
0	0.550781	0.570312
1	0.230835	0.197021
2	0.157837	0.197021
3	0.230835	0.197021
4	0.157837	0.197021

	subtype_code_target_enc
0	0.511230
1	0.230835
2	0.158691
3	0.230835

4 0.158691

[5 rows x 29 columns]

```
[74]: tfidf_reduced_df = pd.concat([pd.DataFrame(X_train), pd.DataFrame(X_test)],  
    ↪axis=0)  
tfidf_reduced_df.columns = ['tfidf_interaction_1', 'tfidf_interaction_2',  
    ↪'tfidf_interaction_3',  
    ↪'tfidf_interaction_4', 'tfidf_interaction_5',  
    ↪'tfidf_interaction_6',  
    ↪'tfidf_interaction_7', 'tfidf_interaction_8',  
    ↪'tfidf_interaction_9',  
    ↪'tfidf_interaction_10']  
print(tfidf_reduced_df.shape)  
tfidf_reduced_df.head()
```

(6639294, 10)

```
[74]: tfidf_interaction_1  tfidf_interaction_2  tfidf_interaction_3  \  
0          0.002103          0.0          0.000000  
1          0.001955          0.0          0.000000  
2          0.001766          0.0          0.035762  
3          0.001955          0.0          0.000000  
4          0.001766          0.0          0.035762  
  
tfidf_interaction_4  tfidf_interaction_5  tfidf_interaction_6  \  
0          0.0          0.0          0.030325  
1          0.0          0.0          0.029987  
2          0.0          0.0          0.030042  
3          0.0          0.0          0.029987  
4          0.0          0.0          0.030042  
  
tfidf_interaction_7  tfidf_interaction_8  tfidf_interaction_9  \  
0          0.015853          0.0          0.0  
1          0.000000          0.0          0.0  
2          0.000000          0.0          0.0  
3          0.000000          0.0          0.0  
4          0.000000          0.0          0.0  
  
tfidf_interaction_10  
0          0.000039  
1          0.000000  
2          0.000000  
3          0.000000  
4          0.000000
```

```
[75]: for col in tfidf_reduced_df.columns:
      print(col)
      test1 = tfidf_reduced_df[col].values
      df1[col] = test1
df1.head()
```

```
tfidf_interaction_1
tfidf_interaction_2
tfidf_interaction_3
tfidf_interaction_4
tfidf_interaction_5
tfidf_interaction_6
tfidf_interaction_7
tfidf_interaction_8
tfidf_interaction_9
tfidf_interaction_10
```

```
[75]:  date_block_num  item_cnt_month  date_avg_item_cnt_lag_1  \
0           12           0.0           0.411133
1           12           0.0           0.411133
2           12           0.0           0.411133
3           12           1.0           0.411133
4           12           1.0           0.411133

      date_item_avg_item_cnt_lag_1  date_item_avg_item_cnt_lag_2  \
0           0.086975           0.044434
1           1.021484           1.022461
2           0.543457           0.600098
3           1.934570           1.799805
4           0.913086           0.333252

      date_item_avg_item_cnt_lag_3  date_item_avg_item_cnt_lag_6  \
0           0.130493           0.065247
1           0.521973           0.891113
2           0.543457           0.304443
3           1.260742           1.891602
4           0.717285           1.000000

      date_item_avg_item_cnt_lag_12  date_shop_avg_item_cnt_lag_1  \
0           0.155518           0.148071
1           0.000000           0.148071
2           0.000000           0.148071
3           5.378906           0.148071
4           1.355469           0.148071

      date_shop_avg_item_cnt_lag_2  ...  tfidf_interaction_1  \
0           0.100647  ...           0.002103
1           0.100647  ...           0.001955
```

2	0.100647	...	0.001766
3	0.100647	...	0.001955
4	0.100647	...	0.001766

	tfidf_interaction_2	tfidf_interaction_3	tfidf_interaction_4	\
0	0.0	0.000000	0.0	
1	0.0	0.000000	0.0	
2	0.0	0.035762	0.0	
3	0.0	0.000000	0.0	
4	0.0	0.035762	0.0	

	tfidf_interaction_5	tfidf_interaction_6	tfidf_interaction_7	\
0	0.0	0.030325	0.015853	
1	0.0	0.029987	0.000000	
2	0.0	0.030042	0.000000	
3	0.0	0.029987	0.000000	
4	0.0	0.030042	0.000000	

	tfidf_interaction_8	tfidf_interaction_9	tfidf_interaction_10
0	0.0	0.0	0.000039
1	0.0	0.0	0.000000
2	0.0	0.0	0.000000
3	0.0	0.0	0.000000
4	0.0	0.0	0.000000

[5 rows x 39 columns]

```
[76]: print(df1.shape)
      print(df1.columns.values)
```

```
(6639294, 39)
['date_block_num' 'item_cnt_month' 'date_avg_item_cnt_lag_1'
 'date_item_avg_item_cnt_lag_1' 'date_item_avg_item_cnt_lag_2'
 'date_item_avg_item_cnt_lag_3' 'date_item_avg_item_cnt_lag_6'
 'date_item_avg_item_cnt_lag_12' 'date_shop_avg_item_cnt_lag_1'
 'date_shop_avg_item_cnt_lag_2' 'date_shop_avg_item_cnt_lag_3'
 'date_shop_avg_item_cnt_lag_6' 'date_shop_avg_item_cnt_lag_12'
 'date_cat_avg_item_cnt_lag_1' 'date_shop_cat_avg_item_cnt_lag_1'
 'date_shop_type_avg_item_cnt_lag_1' 'date_type_avg_item_cnt_lag_1'
 'date_city_avg_item_cnt_lag_1' 'date_shop_subtype_avg_item_cnt_lag_1'
 'date_item_city_avg_item_cnt_lag_1' 'date_subtype_avg_item_cnt_lag_1'
 'delta_price_lag' 'delta_revenue_lag_1' 'shop_id_target_enc'
 'item_id_target_enc' 'city_code_target_enc'
 'item_category_id_x_target_enc' 'type_code_target_enc'
 'subtype_code_target_enc' 'tfidf_interaction_1' 'tfidf_interaction_2'
 'tfidf_interaction_3' 'tfidf_interaction_4' 'tfidf_interaction_5'
 'tfidf_interaction_6' 'tfidf_interaction_7' 'tfidf_interaction_8'
 'tfidf_interaction_9' 'tfidf_interaction_10']
```

```
[77]: df1.to_csv('./data_processed/data.csv') #2.27G
```

```
[ ]: # df1 = pd.read_csv('data.csv', index_col=0)
```